

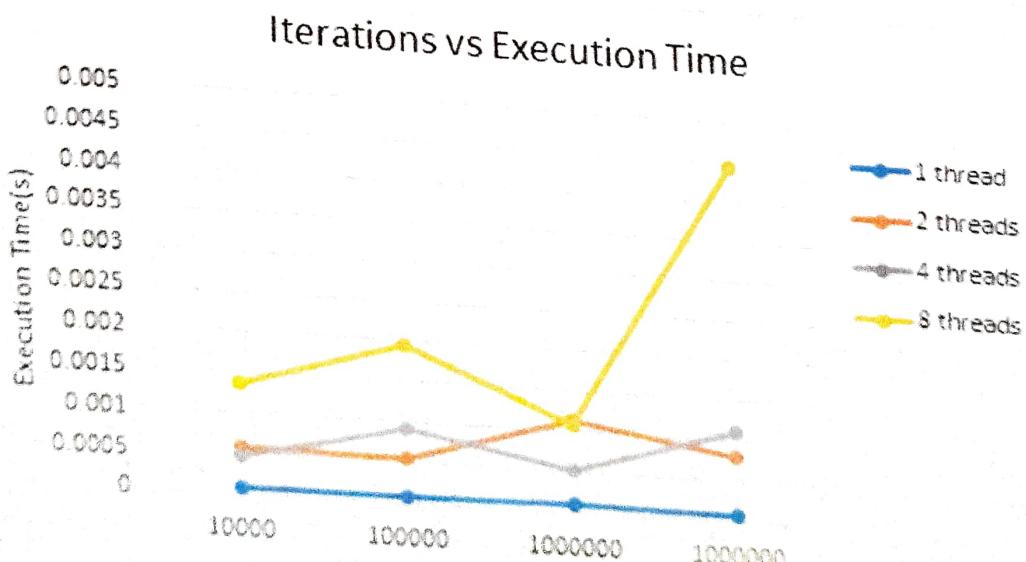
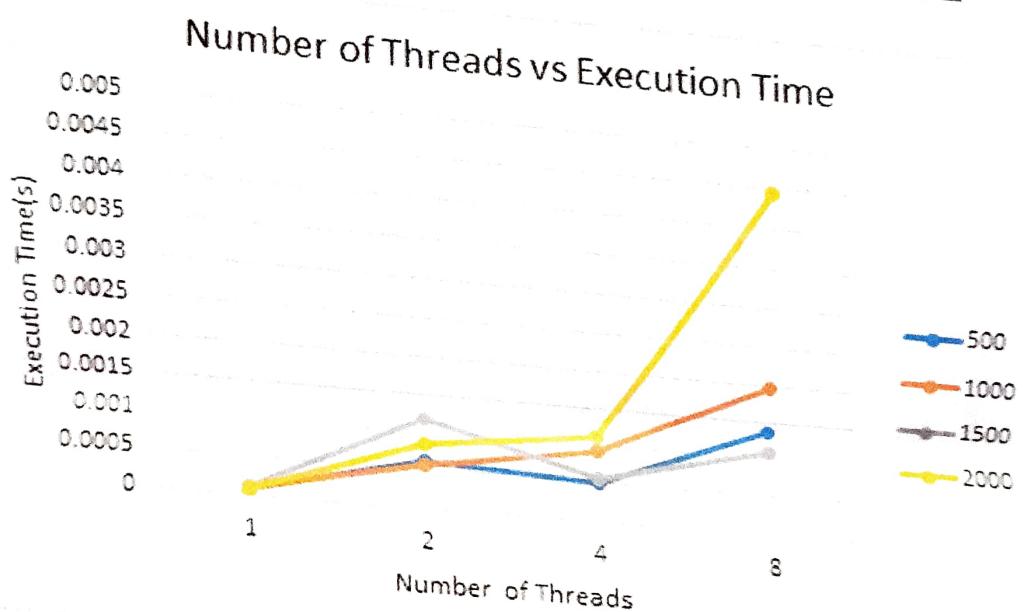
INDEX

St. No.	Experiment	Page	Date of Experiment	Date of Submission	Remark
1.	Monte Carlo Algorithm	1			
2.	A single matrix-matrix multiplication using dynamic memory allocation.	4			
3.	Cache Unfriendly Sieve of Eratosthenes & Parallel Sieve of Eratosthenes	6			
4.	Convert a color image to black & white image	9			
5.	Write a parallel program for points classification	12			
6.	Write a parallel program for word searching a file	14			
7.	Write a MPI-C program multi-tasking	17			
8.	Integral using a quadrature rule	21			

INDEX

Sr. No.	Experiment	Page	Date of Experiment	Date of Submission	Remark
9.	Write a MPI-C program which estimates the time it takes to send a vector of N double precision values through each process in a ring.	24			
10.	Simple matrix multiplication using dynamic memory allocation.	27			
11.	Write an OpenACC program to implement 1D Jacobi iteration.	30			

Iterations	Execution Time-	Number of threads		
	1	2	4	8
10000	0.000068	0.000599		
100000	0.000095	0.000549	0.00048	0.001361
1000000	0.000112	0.001156	0.000909	0.001967
10000000	0.000069	0.000819	0.000543	0.001086
			0.001117	0.004463



AIM

To write an OpenMP program that computes value of PI using Monte Carlo algorithm.

Program IA

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#define SEED 50123
```

```
int main()
```

```
{ long n=0, i, count = 0;
```

```
double x,y,z;
```

```
srand (SEED);
```

```
perlf ("size 1+T1 1+T2 1+T4 1+T8");
```

```
for( n=100; n <= 1000000; n += 10)
```

```
perlf ("n 1d1t", n);
```

```
for( int t=1; t<=8; t += 2) {
```

```
count = 0;
```

```
double start = omp_get_wtime();
```

```
#pragma omp parallel for private (x,y,z)
```

```
reduction (+:count) num_threads (t)
```

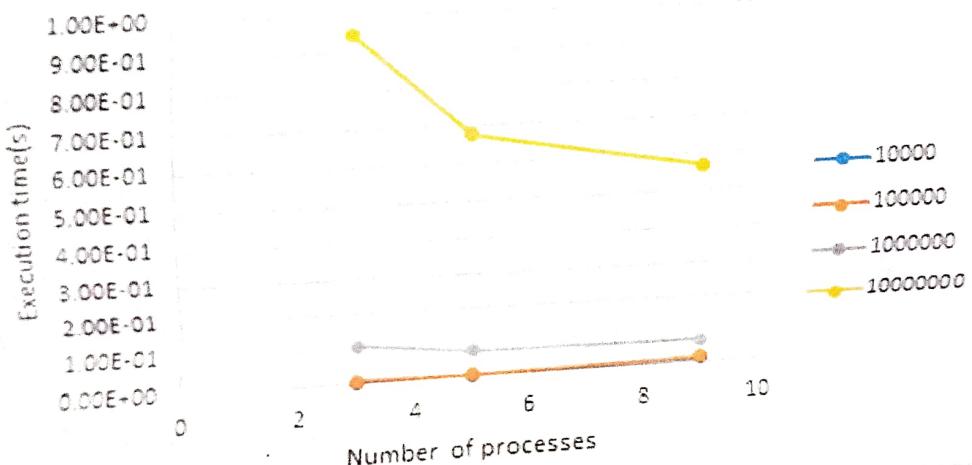
```
for (i=0; i<n; i++)
```

```
{ x = (double) rand () / RAND_MAX;
```

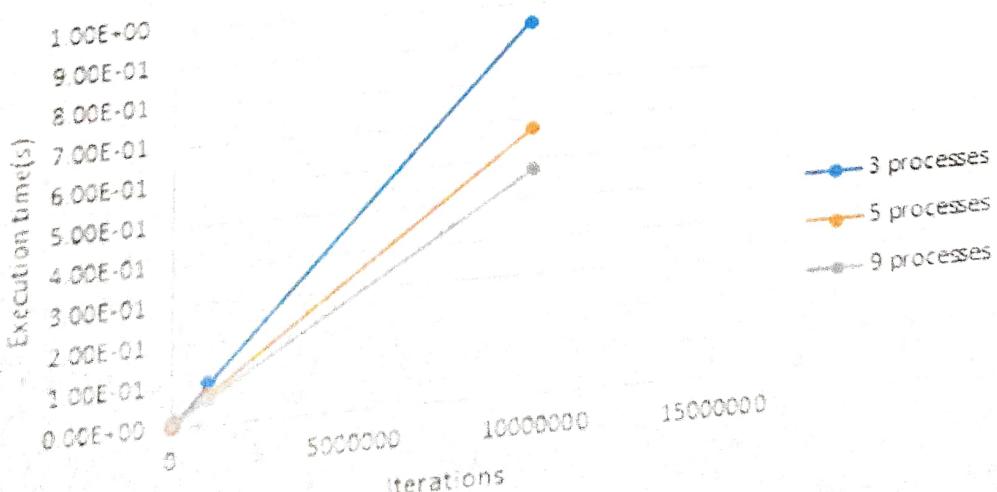
MPI - MONTE CARLO

Iterations	Execution Time		
	3	5	9
10000	1.48E-03	8.57E-04	1.36E-02
100000	9.86E-03	7.58E-03	1.71E-02
1000000	0.105168	0.074548	0.065742
10000000	0.949605	0.67502	0.567609

Number of Processes vs Execution Time



Iterations vs Execution Time



Experiment No. 1 Name: ADITYA PATHAK

LCS / RAND

count / n

t_{uniform}
 t_{rand}

$t_{\text{pi}, \text{st}}$

μ_m
 σ_C

250

```

y = (double) rand () / RAND_MAX;
z = x*x + y*y;
if (z <= 1) count++;
}
double pi = (double) count / n * 4;
double stop =omp_get_wtime();
printf ("If %fs it", pi, stop-start);
}
printf ("\n");
return 0;
}

```

Program 1B

AIM

Write a MPI Program that computes the value of PI using Monte-Carlo Algorithm

Program 1B

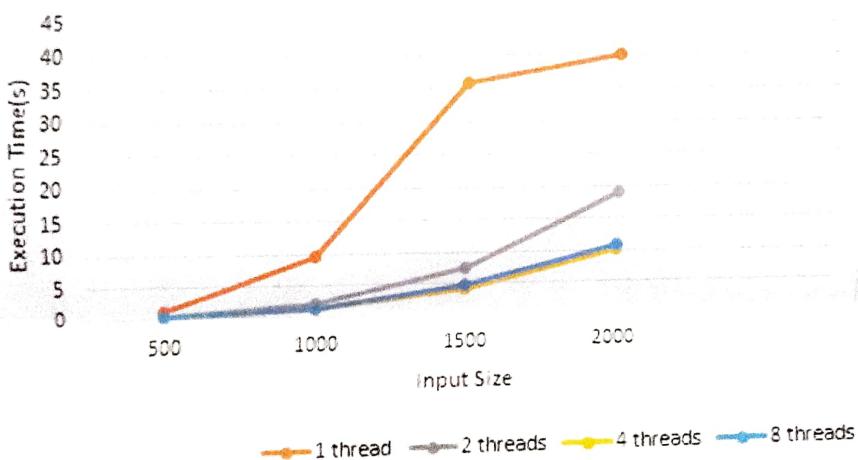
```

#include <stdio.h>
#include <mpi.h>
#include <stdlib.h>
#define SEED 3655942

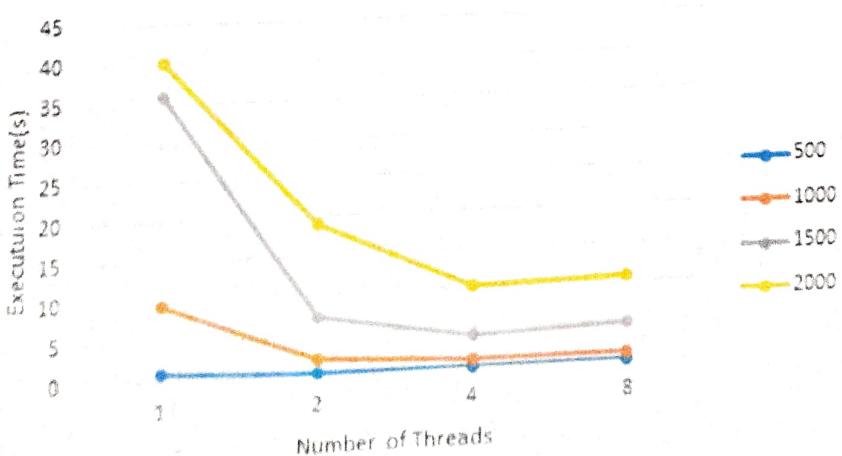
```

Input size	Execution Time (1 thread)	Execution Time (2 threads)	Execution Time (4 threads)	Execution Time (8 threads)
500	1.047931	0.219803	0.12051	0.126442
1000	9.559035	1.972103	1.072029	1.1002
1500	35.79007	7.371719	4.025677	4.717746
2000	40.05189	19.12589	10.25357	10.71961

Input Size vs Execution Time



Number of Threads vs Execution Time



Experiment No. 2
Name:

ADITYA PATI

Date:

1. $\text{int}^{\star} \text{size}$
2. $(\text{int}^{\star} \text{size})$
3. $\text{int}^{\star} \text{size}$
4. $\text{int}^{\star} \text{size}$

5. $\text{int}^{\star} \text{size}$
6. $\text{int}^{\star} \text{size}$
7. $\text{int}^{\star} \text{size}$

in thread computation.
Efficiency is
increased when
we use more
than one thread
and using the
lock mechanism.

AIM

Write an OpenMP program that computes a simple matrix-matrix multiplication using dynamic memory allocation.

- a) Illustrate the correctness of the program
- b) Justify the inference when outer 'for loop' is parallelized with & without using the explicit data scope variables.

Program

```
#include <stdio.h>
#include <omp.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    printf ("Size 1t, 1t\1t\1t\n");
    for (int n=500; n<=2000; n+=500)
    {
        printf ("n\tcl 1t", n);
        1t *ta, *b, *c;
        int i, j, k;
        a = (int **) malloc (n * sizeof (int *));
        b = (int **) malloc (n * sizeof (int *));
        c = (int **) malloc (n * sizeof (int *));
        for (i=0; i<n; i++)
            a[i] = (int *) malloc (n * sizeof (int));
    }
}
```

$b[i] = (\text{int} *) \text{malloc} (n * \text{sizeof} (\text{int}))$;

$c[i] = (\text{int} *) \text{malloc} (n * \text{sizeof} (\text{int}))$;

for ($i=0; i < n; i++$) {

$a[i] = [j] = \text{rand} () \cdot n$;

$b[i][j] = \text{rand} () \cdot n$;

}}

for ($\text{int } t=1; t \leq 8; t += 2$) {

double x = omp_get_wtime();

pragma omp parallel for private (j, k) num_threads (t)

for ($i=0; i < n; i++$) {

for ($j=0; j < n; j++$) {

$c[i][j] = 0$;

for ($k=0; k < n; k++$) {

$c[i][j] += a[i][k] * b[k][i];$

double y = omp_get_wtime();

printf ("%.0f It %f - %f) i }}

printf ("\n");

return 0;

}

AIM

Write a program for Cache unfriendly Sieve of Eratosthenes, Cache friendly Sieve of Eratosthenes & Parallel Sieve for enumerating numbers upto N & prove the correctness.

Program

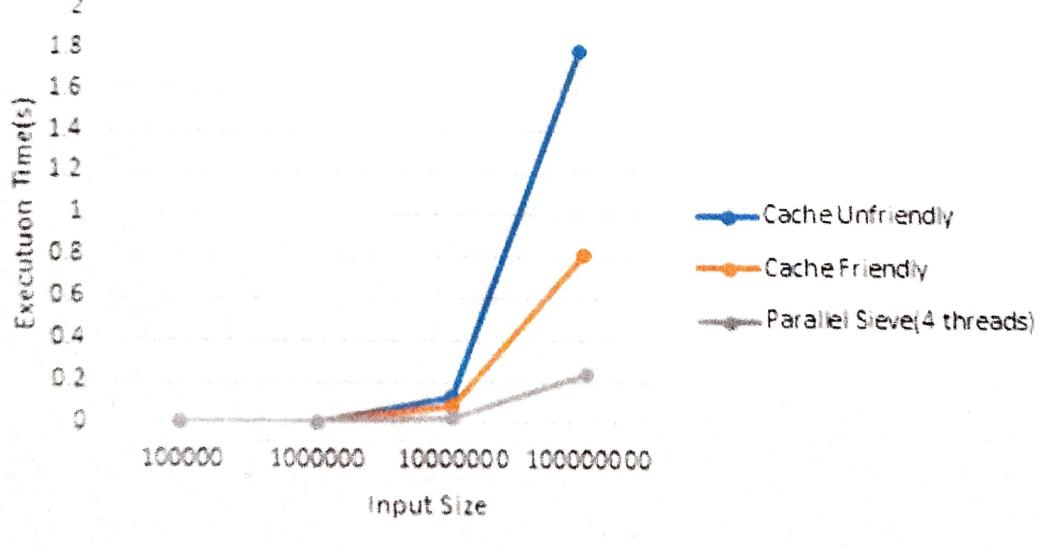
```
#include <math.h>
#include <string.h>
#include <omp.h>
#include <iostream>
using namespace std;
double t = 0;

long strike ( bool composite[], long i, long stride, long limit ) {
    for ( ; i < limit; i += stride )
        composite[i] = true;
    return i;
}
```

```
long Cache Unfriendly Sieve ( long n ) {
    long count = 0;
    m = ( long ) sqrt ( ( double ) n );
    bool * composit = new bool[n+1];
    memset ( composit, 0, n );
    t =omp_set_wtime();
}
```

Input size	Execution Time						
	Cache Unfriendly Sieve	Cache Friendly Sieve	Cache Parallel Sieve	1	2	4	8
100000	0.000818	0.000913	0.00103	0.001074	0.000896	0.001647	
1000000	0.008269	0.009027	0.008996	0.005949	0.003532	0.004876	
10000000	0.122772	0.079445	0.082135	0.052763	0.025498	0.019586	
1E+08	1.79078	0.805944	0.797903	0.422097	0.245255	0.245255	

Sieve Of Erathostenes



```

for (long i=2; i<=m; ++i)
    if (!composite[i])
        ++count;
        strike(composite, 2*i, i, n);
}

```

```

for (long i=m+1; i<=n; ++i)
    if (!composite[i]) ++count;
    t =omp_get_wtime() - t;
    delete [] composite;
    return count;
}

```

`long Parallel Sieve (long n) {`
`long count = 0; m = (long)sqrt((double)n);`
`n-factor = 0;`
`long * factor = new long [m];`
`t = omp_get_wtime();`

`#pragma omp parallel`
`{`

`bool * composite = new bool [m+1]; long * strike`
`= new long [m];`

`#pragma omp single`
`{`

`memset(composite, 0, m);`
`for (long i=2; i<=m; ++i)`

```

if (!composite, 2 * i, m)
    factor[n-factor + f] = i
}

```

long base = -1;

program opr for reduction (+:count)

```

for (long window = m + 1; window <= n; window += m)

```

```

    memset (composite, 0, m);

```

```

    if (base != window) {

```

```

        base = window;
    }
}

```

int main() {

```

    long size = 10000, count;

```

```

    printf ("Size %ld Cache Unfriendly %ld Cache
friendly %ld Parallel Sieve (%ld");

```

```

    for (int i = 1; i <= 4; i++) {
        size = size * 10;
    }

```

```

    printf ("%ld %ld", size);

```

```

    count = Cache_Unfriendly_Sieve (size);

```

```

    printf ("%ld %ld %f %f", count, t);

```

```

    count = Parallel_Sieve (size);

```

```

    printf ("%ld %ld %f %f", count, t);
}

```

```

    return 0;
}

```

ADITYA PATTI

Q1

Write a program to convert a colour image to black & white image.

a) Generate the performance of different scheduling techniques for varying chunk values.

b) Analyze the scheduling patterns by assigning a single color value for an image for each thread.

Program

```
#include < stdio.h >
#include < error.h >
#include < gd.h >
#include < string.h >
#include < config.h >
```

```
int main(int argc, char ** argv) {
    FILE * fp, * fpi = fopen(argv[1], "r");
    gdImagePtr im;
    char name[16], fname[15];
    int color, r, g, b, i = 0, red, green, blue, transparency;
    long width, height;
    config_sched_t df_sched; int def_chunk_size;
    config_gd_sched_t (df_sched, def_chunk_size);
```

Scheduling	Type		
Default	Static	Dynamic	Guided
0.00097	0.000713	0.000824	0.000786
0.001888	0.001904	0.002056	0.001899
0.00676	0.006501	0.007738	0.006864
0.025445	0.025573	0.025754	0.020828

Scheduling Type	Chunk Size	Dynamic	Guided
	50	0.000824	0.001588
	100	0.001875	0.001884
	150	0.001751	0.001774
	200	0.002351	0.002932

Chunk Size vs Execution Time

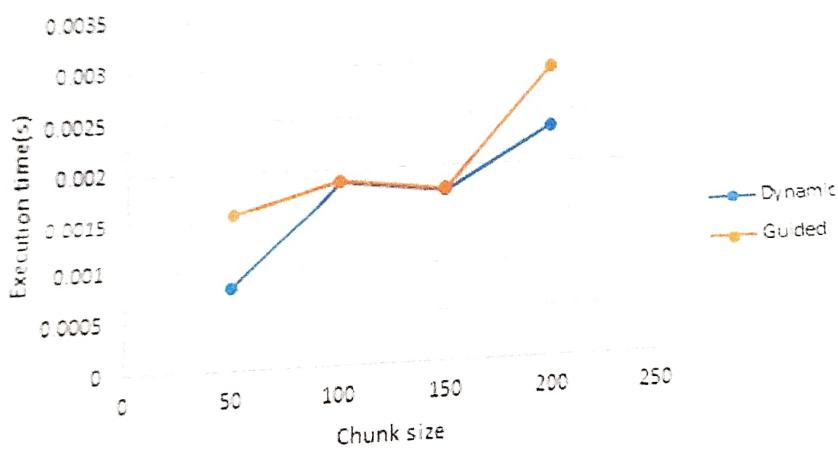
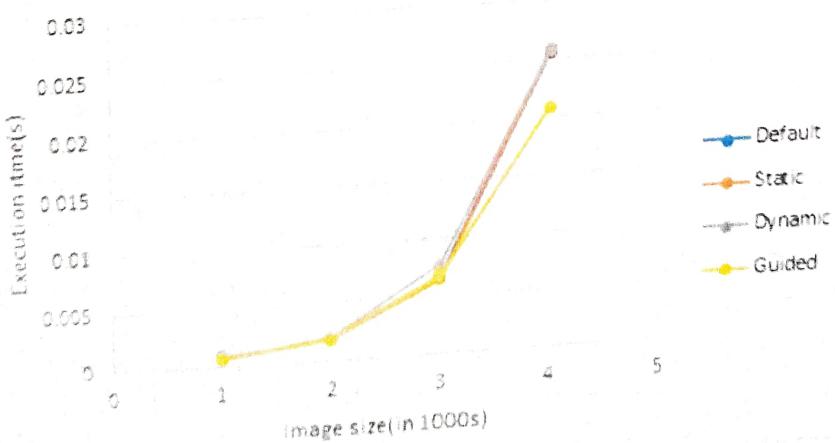


Image Size vs Execution Time



PATI
convert a c
image.
choice of di
for various
ing patterns
for an image

```

for (int i=0; i<4; i++) {
    sprintf (name, "%d%d.png", i+1);
    for (int sched = 0x0; sched <= 0x3; sched++) {
        fp = fopen(name, "r");
        sprintf (name, "Output%d%d.png", i+1, sched);
        img = gdImageCreatefromPng(fp);
        w = gdImageSX(img);
        h = gdImageSY(img);
        if (sched == 0x0)
            printf ("%d%d %d %d\n", w, h);
        omp_set_schedule(def_sched, def_chunk_size);
    }
}
else
omp_set_schedule(sched, 0);
double t = omp_get_wtime();
#pragma omp parallel for private(y, color, red,
green, blue, tmp, tid)
for (x=0; x<w; x++) {
    for (y=0; y<h; y++) {
        color = gdImageGetPixel(img, x, y);
        red = gdImageRed(img, color);
        green = gdImageGreen(img, color);
        blue = gdImageBlue(img, color);
        red = green = blue = tmp;
        gdImageSetPixel(img, x, y, color);
    }
}

```

```
t = omp_get_wtime() - t;  
fp1 = fopen (oname, "w");  
gdImagePng (img, fp1);  
fclose (fp1);  
gdImageDestroy (img);  
rewif ("%.6f\n", t); }  
rewif ("\n"); } return 0; }
```

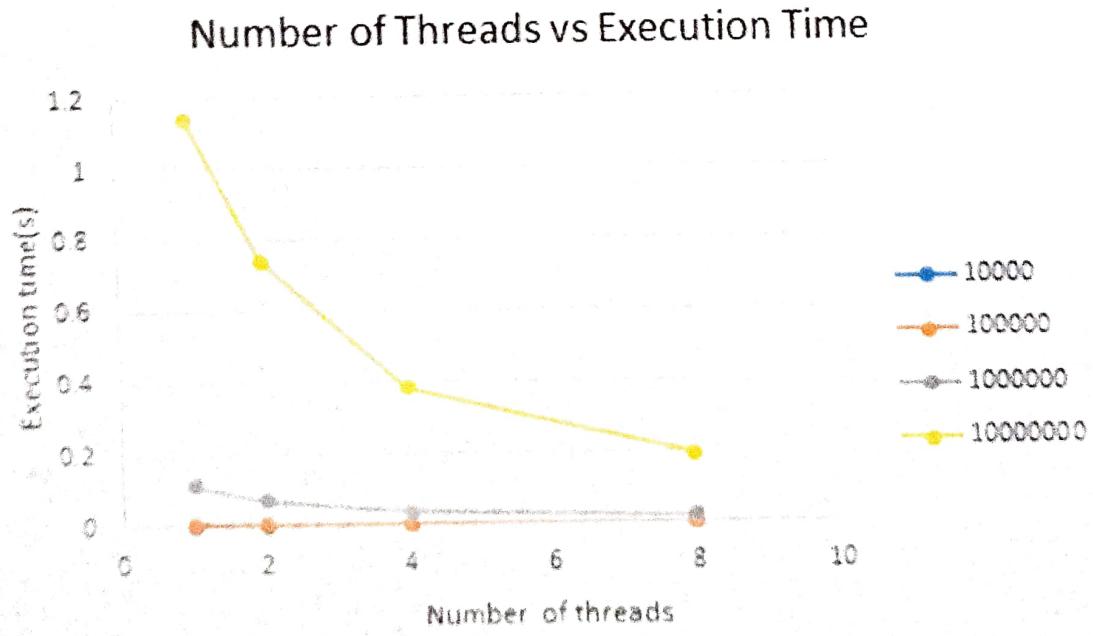
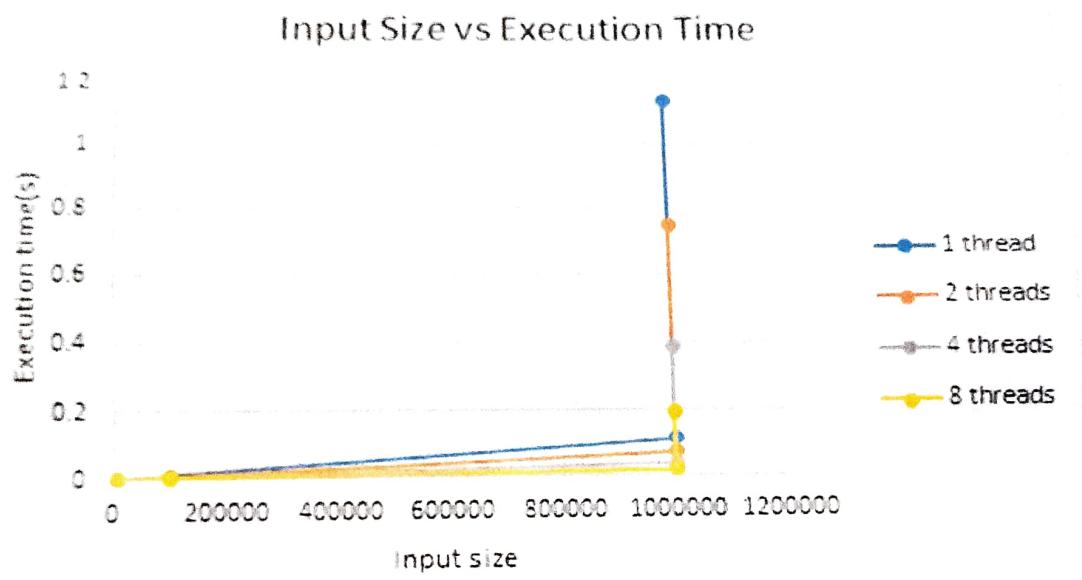
AIM

Write a parallel program for Points Classification. Justify the inference.

Program

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <math.h>
#define CLUSTER_SIZE 4
int cluster [CLUSTER_SIZE][2] = {{75, 25}, {25,
25}, {25, 75}, {75, 75}};
long long cluster_count [CLUSTER_SIZE];
short points [10000000][2];
char output [10000] = " ";
void populate_points (unsigned long long size) {
    long long i;
    for (i=0; i<size; i++) {
        srand (i);
        points [i][0] = rand () % 100;
        points [i][1] = rand () % 100;
    }
}
double get_distance (int x1, int y1, int x2, int y2) {
    int x = x2 - x1, y = y2 - y1;
}
```

Input Size	Execution Time			
	1	2	4	8
10000	0.001208	0.001096	0.000922	0.00122
100000	0.011513	0.00773	0.00442	0.002891
1000000	0.114586	0.074563	0.039273	0.020269
10000000	1.141697	0.742433	0.383055	0.193872



```
return (double) sqrt ((x*x) + (y*y));
```

```
}
```

```
int main (SE
```

```
double t;
```

```
printf (output, "Size | T1 | T2 | T3 | n");
```

```
for (int index = 0; index < 5; index++) {
```

```
printf ("Size: %ld", points_sizes [index]);
```

```
unsigned long long i;
```

```
for (int nt = 1; nt < 9; nt += 2) {
```

```
for (i = 0; i < CLUSTER_SIZE; i++)
```

```
cluster_count [i] = 0;
```

```
t = ovp_get_wtime();
```

```
# pragma omp parallel for reduction (+:cluster_out)
```

```
num_threads (nt)
```

```
for (i = 0; i < points_sizes [index]; i++) {
```

```
double min_dist = 100, cur_dist = -1;
```

```
int j, cluster_index = -1;
```

```
for (j = 0; j < CLUSTER_SIZE; j++) {
```

```
cur_dist = get_distance (points [i] [0],
```

```
points [i] [1], cluster [j] [0], cluster [j] [1]);
```

```
if (cur_dist < min_dist) {
```

```
min_dist = cur_dist;
```

```
cluster_index = j; } }
```

```
puts (output);
```

```
return 0; }
```

AIM

Write a parallel program for Word Search in a file. Justify the inference.

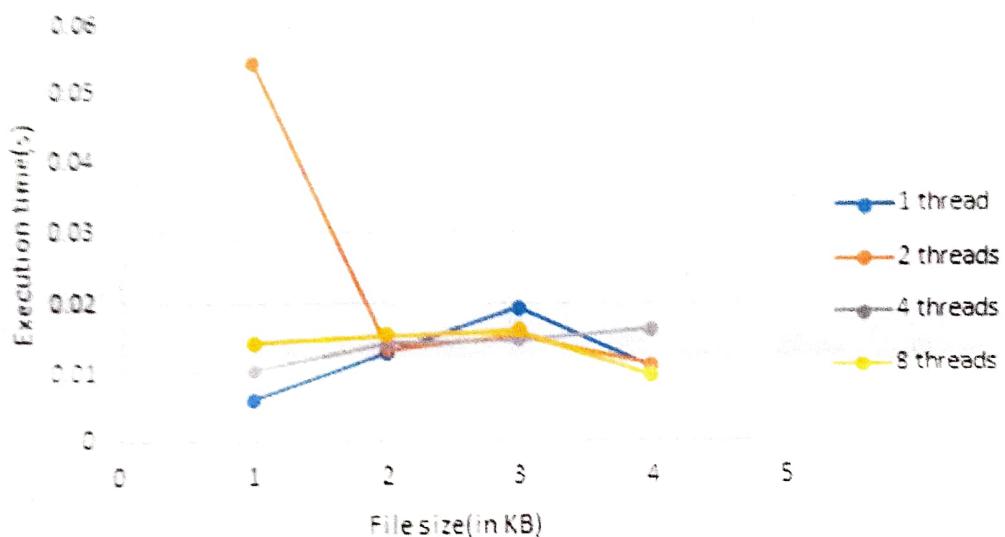
Program

```
#include <stdio.h>
#include <omp.h>
#include <ctype.h>
#include <string.h>
#define COUNT 10
char search_words[20][COUNT] = { "The", "around",
                                "graphics", "from", "by", "be", "a", "which", "various",
                                "mount" };
long counts[COUNT];
int is_equal(char*a, const char*key, int ignore_case) {
    int len_a = strlen(a), len_b = strlen(key);
    if (len_a != len_b) return 0;
    if (ignore_case) returnstrcasecmp(a, key) == 0;
    return strcmp(a, key) == 0;
}
void read_word(char *temp, FILE *fp) {
    int i = 0; char ch;
```

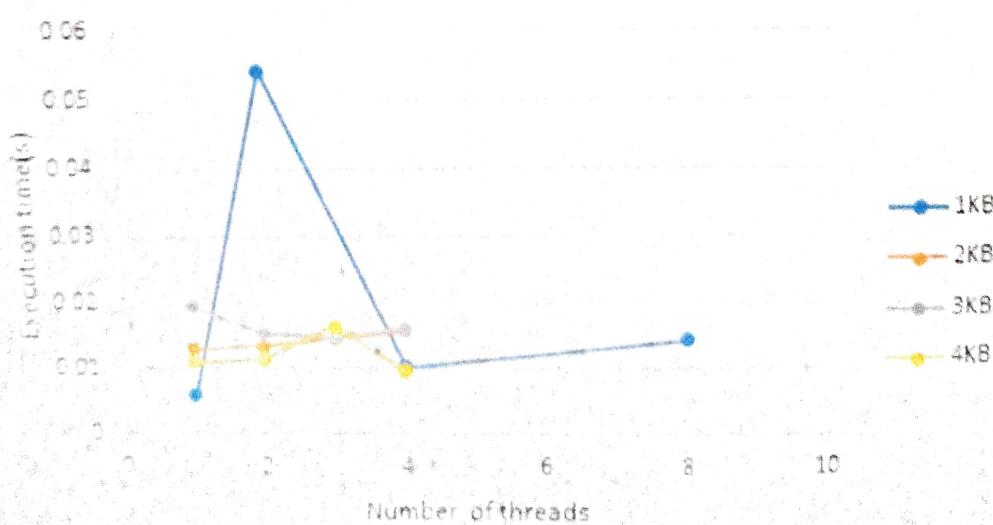
Execution Time

File Size(KB)	1	2	4	8
1	0.006009	0.054372	0.010441	0.014479
2	0.013115	0.013621	0.014502	0.015667
3	0.019483	0.015392	0.014927	0.016296
4	0.01105	0.0115	0.016502	0.009686

File Size vs Execution Time



Number of Threads vs Execution Time



while ((ch = fgetc(fp)) != EOF && isalpha(ch))
 == 0) { }

while (ch != EOF && isalpha(ch) != 0) {
 temp[i + t] = ch;
 ch = fgetc(fp); }
 temp[i] = '\0'; }

long determine_count (const char* file-name, const
 char* key, int ignore_case) {
 int key_index = 0, key_len = strlen(key);
 long word_count = 0; char ch;
 FILE *fp = fopen(file-name, "r");
 char temp[50]; int i = 0;
 while (feof(fp) == 0) { read_word(temp, fp);
 if (!is_equal(temp, key, ignore-case) != 0)
 word_count++; }
 return word_count; }

int main()

{ char output[1000] = " ";
 int i; for (i = 0; i < COUNT; i++)
 count[i] = 0;
 char *myfiles[4] = {"file1.txt", "file2.txt",
 "file3.txt", "file4.txt"};
 sprintf(output, "Size %T1 %T2 %T3 %T4"); }

```
for (int ita = 0; ita < 4; ita++) {  
    FILE *fp = fopen("my-file[ite].txt", "w");  
    fseek(fp, 0, SEEK_END);  
    for (int t = 1; t <= 8; t += 2) {  
       omp_set_num_threads(t);  
        double start = omp_get_wtime();  
        #pragma omp parallel for  
        for (int i = 1; i <= 8; i += 2) {  
            omp_set_num_threads(i);  
            fprintf(output, "%s\n", search_words[i],  
                    count[i]);  
        }  
        printf("\n");  
        auto output;  
    }  
}
```

AIM

Write a MPI-C program Multi-Tasking.
Justify the inference.

Program

```
#include <math.h>
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
void po_set_input (int *input1, int *input2) {
    *input1 = 10000000;
    *input2 = 100000;
    printf ("In PO_SET_PARAMETERS", *input1, *input2);
}

void po_send_input (int input1, int input2) {
    int id=1, tag=1;
    MPI_Send (&input1, 1, MPI_INT, id, tag, MPI_COMM_WORLD);
    id=2, tag=2;
    MPI_Send (&input2, 1, MPI_INT, id, tag, MPI_COMM_WORLD);
}
```

```
aditya@aditya-Lenovo-ideapad-320-15IKB:~/Dow
10 January 2021 05:30:01 PM

MPI_MULTITASK:
  C / MPI version

P0_SET_PARAMETERS:
  Set INPUT1 = 10000000
  INPUT2 = 100000
  Process 2 time = 1.47705
  Process 1 time = 7.72909

  Process 1 returned OUTPUT1 = 615
  Process 2 returned OUTPUT2 = 9592
  Process 0 time = 7.72904

MPI_MULTITASK:
  Normal end of execution.
10 January 2021 05:30:08 PM
```

int pl_semicne_input () {

 int id = 0, input1, tag = 1; MPI_Status status;
 MPI_Recv (&input1, 1, MPI_INT, id, tag, MPI_COMM_WORLD,
 &status); return input1; }

int pl_compute_output (int input1) {

 int i, j, k, output1 = 0;
 for (i = 2; i <= output1; i++) {
 j = i; k = 0;
 while (j < i) {
 if ((j % 2) == 0) j = j / 2;
 else j = 3 * j + 1;
 k++;

}

 if (output1 < k) output1 = k; }
 return output1;

}

void pl_send_output (int output1) {

 int id = 0, tag = 2; MPI_Send (&output1, 1, MPI_INT, id, tag,
 MPI_COMM_WORLD);

}

int pl_semicne_input () {

 int id = 0, input2, tag = 2; MPI_Status status;
 MPI_Recv (&input2, 1, MPI_INT, id, tag, MPI_COMM_WORLD,
 &status); return input2; }

```

int p2_compute_output (int input2) {
    int i, j, output2 = 0, prime;
    for (i=2; i<=input2; i++) {
        prime = i;
        for (j=2; j<i; j++) {
            if (i % j == 0) {
                prime = 0; break;
            }
        }
        if (prime != 0)
            output2++;
    }
    return output2;
}

```

```

Void timestamp () {
#define TIME-SIZE 40
    static char time-buffer[TIME-SIZE];
    const struct tm *tm;
    tm = now();
    now = time(NULL);
    tm = localtime(&now);
    printf ("%s\n", time-buffer);
#undef TIME-SIZE
}

```

```

int main (int argc, char **argv) {
    int id, ierr, input1, input2, output1, output2, pi;
    double wtime;
    ierr = MPI_Init (&argc, argv);
    if (ierr != 0)
        exit(1);
    pi = p2_compute_output(input2);
    output1 = timestamp();
    output2 = timestamp();
    wtime = MPI_Wtime();
    MPI_Finalize();
}

```

Experiment No. _____

7

Name: _____

```

    ierr = MPI_Comm_rank ( MPI_COMM_WORLD, &id );
    MPI_Comm_size ( MPI_COMM_WORLD, &P );
    if ( P < 3 ) {
        MPI_Finalize ();
        exit ( 1 );
    } else if ( id == 0 ) {
        timestamp ();
        wtime = MPI_Wtime ();
        p0_get_input (&input1, &input2);
        p0_receive_output (&output1, &output2);
        wtime = MPI_Wtime () - wtime;
        MPI_Finalize ();
        timestamp ();
    } else if ( id == 1 ) {
        wtime = MPI_Wtime ();
        input1 = p1_receive_input ();
        output1 = p1_compute_output (input1);
        p1_send_output (output1);
        MPI_Finalize ();
    } else if ( id == 2 ) {
        wtime = MPI_Wtime ();
        output2 = p2_compute_output (input2);
        p2_send_output (output2);
        wtime = MPI_Wtime () - wtime;
        MPI_Finalize ();
    }

```

AIM

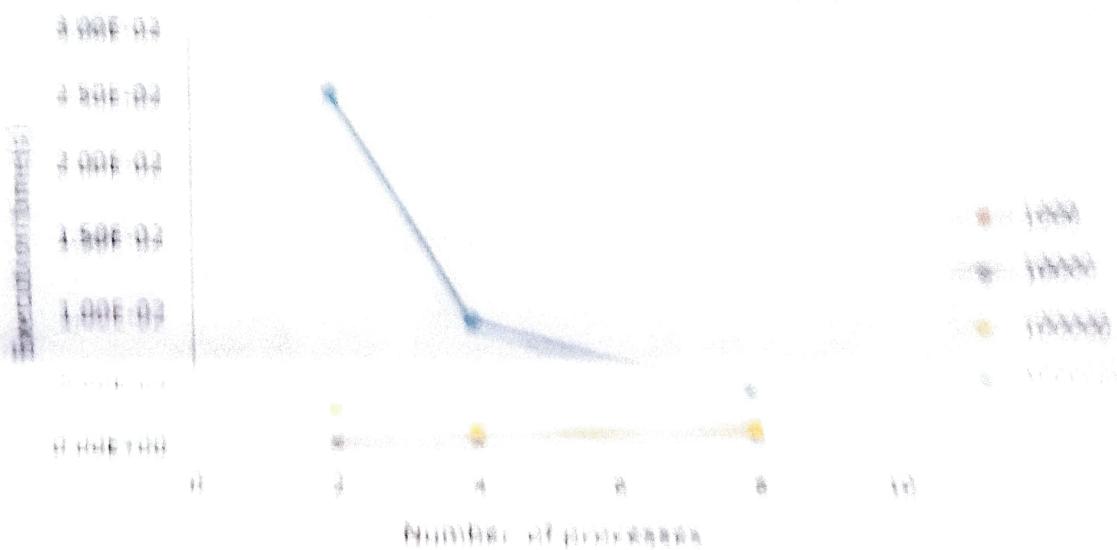
Write MPI-C program which approximates an integral using a quadrature rule.

Program

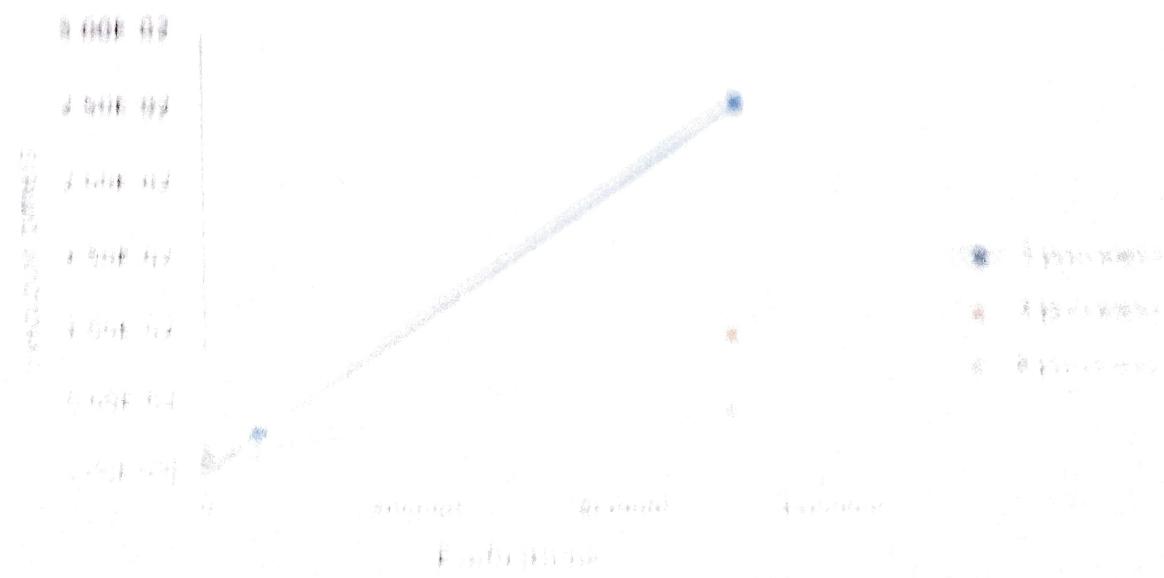
```
#include <math.h>
#include <mpish>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
double f (double x) {
    double pi, value;
    pi = 3.141592653589793;
    value = 50.0 / (pi * (2500 + x * x + 1.0));
    return value;
}
void timestamp(void) {
#define TIME_SIZE 40
    static char time_buffer[TIME_SIZE];
    const struct tm *tm;
    time_t now; now = time(NULL);
    tm = localtime(&now);
    strftime(time_buffer, TIME_SIZE, "%d %B %Y %T %M %S %P", tm);
    printf("%s\n", time_buffer);
    return;
}
```

Evaluations	Execution Time		
	3	8	9
1000	5.50E-04	6.00E-04	1.00E-03
10000	8.00E-04	7.00E-04	9.00E-04
100000	0.003111	0.001600	0.001300
1000000	0.0250	0.009700	0.007000

Number of Processes vs Execution Time



Evaluations vs Execution Time



#include TIME_SIZE }

int main (int argc , char *argv[]) {

double a,b, error, exact, my_total, total, wtime, r;

int i, master = 0, my_id, my_h, h, p, p_num, src;

double my_a, my_b; MPI_Status status;

int tag, target; a = 0; b = 10; n = 10000000;

exact = 0.4993633810764545674464;

MPI_Init (&argc, &argv);

MPI_Comm_Rank (MPI_COMM_WORLD, &my_id);

MPI_Comm_size (MPI_COMM_WORLD, &p_num);

if (my_id == master) {

my_n = n / (p_num - 1);

n = (p_num - 1) # my_n; wtime = MPI_Wtime();

printf ("QUAD MPI - C / MPI version \n");

printf ("A = %f \n", a); printf ("B = %f \n", b);

}

src = master; MPI_Bcast (&my_n, 1, MPI_INT, src, MPI_COMM_WORLD);

if (my_id == master) {

for (p = 1; p <= p_num - 1; p++) {

my_a = ((double) (pnum - p)) * a +

(double) (p - 1) * b / (double) (pnum - 1);

taget = p; tag = 1;

MPI_Send (&my_a, 1, MPI_DOUBLE, target, tag,

MPI_COMM_WORLD);

```

target = p; tag = 2;
MPI_Send (&my_b, 1, MPI_DOUBLE, target, tag,
          MPI_COMM_WORLD); }

total = 0.0; my_total = 0.0; }

else { source = master;
        tag = 1; MPI_Recv (&my_a, 1, MPI_DOUBLE,
                           source, tag, MPI_COMM_WORLD, &status);
        my_total = 0.0;
        for (i=1; i<= my_n; i++) {
            x = ((double) (my_n-i)) * my_a +
                (double) (i-1) * my_b) / double (my_n-1);
            my_total = my_total + f(x); }

        my_total = (my_b-my_a) * my_total / (double) my_n;
        MPI_Reduce (&my_total, &total, 1, MPI_DOUBLE,
                    MPI_SUM, master, MPI_COMM_WORLD);

        if (my_id == master) {
            error = fabs (total - exact);
            wtime = MPI_Wtime () - wtime; }

MPI_Finalize ();
if (my_id == master)
    printf ("End of Execution");
return 0;
}

```

AIM

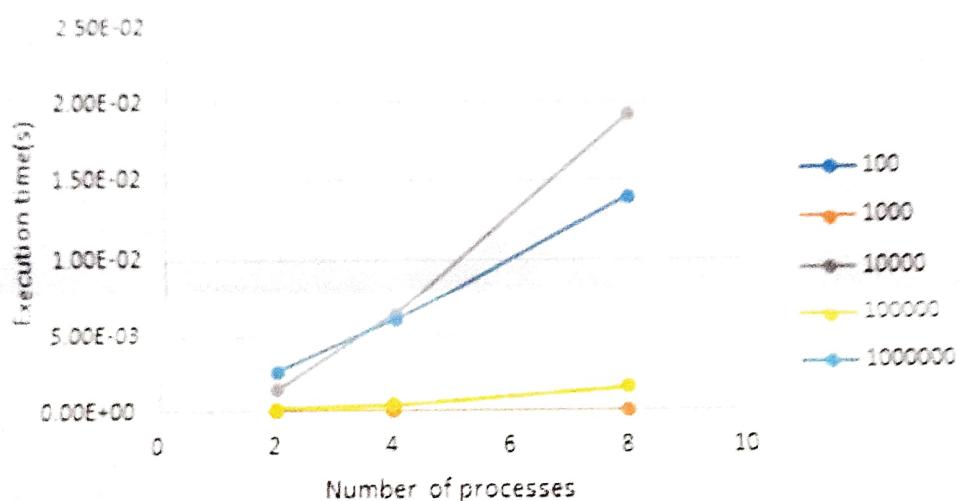
Write a MPI-C program which estimates the time it takes to send a vector of N double precision values through each process in a ring.

Program

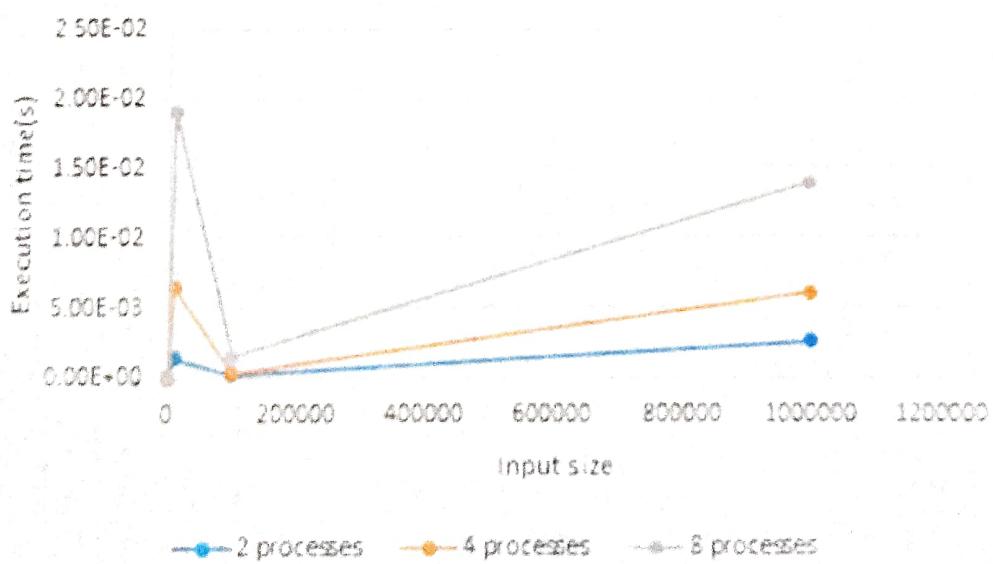
```
#include <mpi.h>
#include <stdlib.h>
#include <stdio.h>
void ring-io (int p, int id)
int main (int argc, char * argv[])
int error, p, id;
MPI_Init (&argc, &argv);
MPI_Comm_size (MPI_COMM_WORLD, &p);
MPI_Comm_rank (MPI_COMM_WORLD, &id);
if (id == 0)
    printf ("In Ring MPI:\n");
    printf ("A ring of processes\n");
    printf ("\n");
    printf ("The no. of processes is %d\n", p);
}
ring-io (p, id);
```

Input Size	Execution Time		
	2	4	8
100	5.27E-05	1.67E-05	6.29E-05
1000	1.07E-05	1.48E-05	6.76E-05
10000	0.001448	0.006482	0.019143
100000	0.000243	0.000422	0.001558
1000000	0.002617	0.006074	0.013918

Number of Processes vs Execution Time



Input Size vs Execution Time



```

MPI_Finalize();
if (id == 0) {
    coutff ("Ring MPI: \n");
}
    
```

Action 0)

```

void ring-io (int p, int id) {
    int dest, i, j, n;
    int n-test[5] = {100, 1000, 10000, 100000};
    int n-test-num = 5, source;
    MPI_Status status; double tave;
    int test, test-num = 10; double tmax, tmin, wtime;
    double *x; if (id == 0) {
        coutff ("Times based on experiment");
        coutff ("tN tT Tave t Tmax \n");
    }
    
```

```

for (i=0; i<n; i++) {
    n = n-test[i];
    x = (double *) malloc (n * sizeof (double));
    if (id == 0) {
        dest = 1; source = p - 1;
        tave = 0.0; tmin = 1.0E+30; tmax = 0.0;
    }
    for (test=1; test <= min; test++) {
        for (j=0; j<n; j++)
            x[j] = (double) (test + j);
        wtime = MPI_Wtime () - wtime;
        tave = tave + wtime;
        if (wtime < tmin) tmin = wtime;
    }
}
    
```

```
if ( tmax < wtime ) tmax = wtime;  
} if ( tave > tave / (double) (test_num); }  
else { source = id - 1;  
dest = ((id + 1) % p);  
for ( test = 1; test <= test_num; test++ ) {  
MPI_Recv (&x, n, MPI_DOUBLE, source, 0, MPI_COMM_WORLD,  
&status);  
MPI_Send (&x, n, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);  
}  
free(x); } return; }
```

AIM

Write an OpenACC program that computes a simple matrix-matrix multiplication using dynamic memory allocation.

Program

```

#include <sys/time.h>
#include <stdlib.h>
#include <stdio.h>
#define MAX 1000
int size;
double **a, **b, **c, **d;
int main (int argc, char *argv[])
{
    SIZE = atoi(argv[1]);
    a = (double **) malloc (SIZE * sizeof (double));
    b = (double **) malloc (SIZE * sizeof (double));
    c = (double **) malloc (SIZE * sizeof (double));
    d = (double **) malloc (SIZE * sizeof (double));
    int i, j, k;
    struct timeval, tim;
    double t1, t2, tnp;
    for (i=0; i<SIZE; ++i) {
        a[i] = (double *) malloc (SIZE * sizeof (double));
        b[i] = (double *) malloc (SIZE * sizeof (double));
        c[i] = (double *) malloc (SIZE * sizeof (double));
    }
}

```

```

for (j=0; j<SIZE; ++j) {
    a[i][j] = (double)(i+j);
    b[i][j] = (double)(i-j);
    c[i][j] = 0.0f;
    d[i][j] = 0.0f; } }

for (i=0; i<SIZE; ++i) {
    for (j=0; j<SIZE; ++j) {
        tmp = 0.0f;
        for (k=0; k<SIZE; ++k)
            tmp += a[i][k] * b[k][j];
        d[i][j] = tmp; } }

```

```

gettimeofday (&tim, NULL);
t1 = tim.tv_sec + (tim.tv_usec / 1000000.0);

#pragma acc data copyin (a,b) copy (c)
#pragma acc kernels
#pragma acc loop tile (1000,1000)
for (i=0; i<SIZE; ++i) {
    for (j=0; j<SIZE; ++j) {
        tmp = 0.0f;
        for (k=0; k<SIZE; ++k)
            tmp += a[i][k] * b[k][j];
        c[i][j] = tmp; } }

```

get time of day (& tim, null);

fl = tim.tv_sec + (tim.tv_nsec / 1000000.0);

for (i=0; i<SIZE; ++i)

 for (j=0; j<SIZE; ++j)

 if (c[i][j] != d[i][j]) {

 exit(1); }

 return 0;

}

AIM

Write an OpenACC program to implement two dimensional Jacobi Iteration.

Program

```
#include <math.h>
#include <string.h>
#include <openacc.h>
#include <sys/time.h>
#include <stdio.h>
#define NN 1024
#define NM 1024
float A[NN][NM], Anew[NN][NM];
int main( int argc, char** argv ) {
    int i, j; const int n = NN; const int NM = m;
    iter_max = 1000; const double tol = 1.0e-6;
    double error = 1.0, t1, t2;
    struct timeval, tm;
    memset (A, 0, n * m * sizeof (float));
    memset (Anew, 0, n * m * sizeof (float));
    for (i = 0; i < n; i++) { A[i][0] = 1.0;
        Anew[i][0] = 1.0; }
    }
```

```
gettimeofday (&tm, NULL);
```

$$t1 = tm.tv_sec + (tm.tv_usec / 1000000.0);$$

```
int iter = 0;
```

```
#pragma acc data copy(1) create (Anew)
```

```
while (error > tol && iter < iter_max){
```

```
error = 0.0;
```

```
#pragma acc parallel loop reduction (max : error)
```

```
for (j=1; j<n-1; j++) {
```

```
for (i=1; i<m-1; i++) {
```

$$Anew[j][i] = 0.25 * A[j][i+1] + A[j][i-1]$$

$$+ A[j-1][i] + A[j+1][i]);$$

$$\text{error} = \max (\text{error}, \text{fabs}(Anew[j][i] - A[j][i]));$$

⋮ ⋮

```
#pragma acc kernels
```

```
for (j=1; j<n-1; j++) {
```

```
for (i=1; i<m-1; i++)
```

$$A[j][i] = Anew[j][i];$$

⋮

```
if (iter % 100 == 0) printf ("%s d,%f ln", iter),
```

```
iter++;
```

```
gettimeofday (&tm, NULL);
```

$$t2 = tm.tv_sec + (tm.tv_usec / 1000000.0);$$

```
return 0;
```